MENTAL HEALTH MONITORING SYSTEM

Phase 1 — Backbone Fine-Tuning

Progress Report

Submitted by: Group 54

Platform: Kaggle Notebooks | Date: June 30, 2026

1. Project Overview

The Mental Health Monitoring System is a multi-phase NLP project designed to automatically detect depression, classify mental health severity, and generate clinically coherent reports from raw text input. The system is structured as a four-layer architecture:

- Layer 1 — Transformer backbone (MentalRoBERTa) for contextual text embeddings
- Layer 2 — Graph Attention Network (GAT) for structural reasoning over token dependency graphs
- Layer 3 — Three-head classifier: emotion recognition, mental health category, and severity regression
- Layer 4 — Large Language Model (Groq) for report generation and crisis alerting

Phase 1 covers Layer 1 exclusively. No GAT layers or classifier heads are involved at this stage. The sole objective is to fine-tune the MentalRoBERTa backbone so it produces embeddings rich in mental health and emotional signal before the graph layer is added in Phase 2.

2. Phase 1 Objective

Fine-tune MentalRoBERTa on two datasets in sequence, producing a checkpoint whose embeddings capture both stress-related social media language and fine-grained emotional nuance. The trained checkpoint serves as the frozen backbone input to the Phase 2 GAT layer.

Training Sequence (Design Decision)

The order of fine-tuning is intentional and follows published best practices for curriculum learning:

- Stage 1 — Dreaddit (binary stress classification): simpler objective, converges fast, teaches the backbone the vocabulary and writing style of mental health Reddit posts
- Stage 2 — GoEmotions (28-class multi-label emotion): harder objective introduced after the backbone is already in the right feature space

Starting with a binary task before a 28-class multi-label task prevents the model from being overwhelmed early and leads to better final representations than training on GoEmotions alone from scratch.

3. Environment and Configuration

Platform	Kaggle Notebooks
Hardware	GPU T4 x2 (NVIDIA)
Framework	PyTorch + HuggingFace Transformers
Base Model	mental/mental-roberta-base (HuggingFace)
Model Type	RoBERTa-base pretrained on mental health social media posts
Max Sequence Length	128 tokens
Access Method	HuggingFace token via Kaggle Secrets (gated model)
Checkpoint Format	HuggingFace SavePretrained (.safetensors + config)

4. Stage 1 — Dreaddit Fine-Tuning

4.1 Dataset

Dataset	Dreaddit — Stress Analysis in Social Media

Source	Kaggle (ruchi798/stress-analysis-in-social-media)
Task	Binary stress classification (stressed = 1, not stressed = 0)
Train split	2,412 samples (85% of original train file)
Validation split	426 samples (15% stratified holdout from train)
Test split	715 samples (official dreaddit-test.csv)

4.2 Training Configuration

Loss Function	Cross Entropy (binary classification)
Epochs	4
Batch Size (train)	16
Batch Size (eval)	32
Learning Rate	2e-5 (AdamW optimizer)
Best Model Selection	Highest F1-macro on validation split
Evaluation Metric	Macro F1-score (not accuracy — class imbalance present)

4.3 Results

Split	F1-macro	Notes
Validation	*Not reported*	*Best checkpoint selected by val F1 (load_best_model_at_end=True)*
Test (held-out)	**0.7954**	*Exceeds published MentalBERT/RoBERTa baselines (range 0.70–0.76) on compara tasks*

4.4 Interpretation

A test Macro F1 of 0.7954 on binary stress classification is a strong result. Published baselines using standard BERT-base and RoBERTa-base on similar Reddit mental health tasks report F1 in the 0.70–0.76 range. The MentalRoBERTa backbone's domain-specific pretraining on mental health social media posts provides a meaningful advantage over generic transformer baselines, and the fine-tuning on Dreaddit successfully transferred that advantage to the stress detection task.

The stratified validation split (15%) ensures that class distribution is preserved during training, preventing the optimiser from being misled by imbalanced batches. Macro F1 is used throughout instead of accuracy because accuracy is misleading when class sizes differ — a model that always predicts the majority class can achieve deceptively high accuracy on imbalanced datasets.

5. Stage 2 — GoEmotions Fine-Tuning

5.1 Dataset

Dataset	GoEmotions (Google Research)
Source	Raw CSV release (3 files: goemotions_1/2/3.csv)
Task	Multi-label emotion classification (28 classes including neutral)
Raw rows	211,225 (one row per annotator rating)
After deduplication	58,009 unique comments (standard reported size)
Aggregation strategy	max() across raters — emotion present if any rater flagged it
Train split	49,307 samples (85%)
Validation split	4,351 samples (7.5%)
Test split	4,351 samples (7.5%)

5.2 Training Configuration

Loss Function	BCEWithLogitsLoss (binary cross-entropy per emotion label)
Head type	Multi-label sigmoid (28 outputs, independent per label)
Loaded from	Dreaddit checkpoint (not fresh MentalRoBERTa)
Head replacement	2-class Dreaddit head replaced with fresh 28-class head
Epochs	3
Batch Size (train)	16
Learning Rate	2e-5
Prediction threshold	0.3 sigmoid probability (not 0.5 — GoEmotions labels are sparse)
Evaluation Metric	Macro F1-score across all 28 labels

5.3 Training Progress

Epoch	Avg Train Loss	Val F1-macro	Notes
1	0.1326	**0.4405**	*Best checkpoint saved*
2	0.1144	**0.4667**	*Best checkpoint saved*
3	0.1067	**0.4796**	*Best checkpoint saved ← used as final*

5.4 Test Set Results

GoEmotions Test Macro F1	0.4876
Comparison benchmark	Original GoEmotions paper (BERT-base fine-tuned): 0.46–0.54
Status vs benchmark	Within published range — model is performing correctly

5.5 Interpretation

A test Macro F1 of 0.4876 on 28-class multi-label emotion classification is consistent with what the original GoEmotions paper reports for BERT-based models (0.46–0.54). This is a substantially harder task than binary stress classification for three reasons:

- 28 emotion classes must be predicted independently rather than as a single decision
- Many emotions co-occur in the same text, making the classification boundaries ambiguous
- Class imbalance is severe — rare emotions like 'grief' (49 test samples) and 'pride' (99 test samples) drag down the macro average even when common emotions are well-predicted

The declining training loss across all 3 epochs (0.1326 → 0.1144 → 0.1067) with a correspondingly rising validation F1 confirms stable learning without overfitting. The model is not memorising the training set — it is generalising.

The 0.3 sigmoid threshold (rather than the standard 0.5) is appropriate here because GoEmotions labels are sparse. If the threshold were kept at 0.5, the model would almost never predict rare emotions, collapsing recall for those classes to near zero. At 0.3, the model is more willing to predict present emotions, improving recall without a significant precision penalty on high-frequency classes.

6. Key Outputs and Deliverables

Deliverable	Location	Purpose
`mentalroberta_dreaddit/`	/kaggle/working/	Intermediate Dreaddit checkpoint
`mentalroberta_phase1_final/`	/kaggle/working/ (committed)	Final Phase 1 checkpoint → input to Phase 2

Both checkpoints are saved in HuggingFace format (.safetensors model weights + tokenizer_config.json + tokenizer.json) and are committed as a permanent Kaggle Output version, persisting independently of session timeouts.

7. Phase 1 Summary

Metric	Value	Notes
Dreaddit Test F1-macro	**0.7954**	*Above published RoBERTa baselines (0.70–0.76)*
GoEmotions Val F1-macro (best)	**0.4796**	*Epoch 3 — consistent improvement each epoch*
GoEmotions Test F1-macro	**0.4876**	*Within original GoEmotions paper benchmark range*
GoEmotions training loss (epoch 3)	**0.1067**	*Steady decrease — no overfitting observed*
Total training samples	**~51,700**	*Dreaddit 2,412 + GoEmotions 49,307*
Backbone parameters	**~125M**	*Standard RoBERTa-base architecture*
Final checkpoint size	**~499 MB**	*.safetensors format*

8. Next Steps — Phase 2

Phase 2 takes the Phase 1 checkpoint as input with all transformer weights frozen and trains a Graph Attention Network (GAT) on top of the fixed embeddings. The GAT learns which grammatical relationships between tokens (negation, symptom modifiers, subject-verb) carry the strongest clinical signal.

- Load mentalroberta_phase1_final and freeze all backbone parameters
- Generate spaCy dependency parse graphs for each input text
- Train GATConv layers using PRIMATE 2024 dataset with DSM-5 symptom classification as supervision
- Verify clinical validity: negation edges should show higher average attention weight than determiner edges after convergence
- Save GAT checkpoint as input to Phase 3 three-head training

DAIC-WOZ data use agreement should be submitted immediately if not already done — Phase 3 severity head training requires this dataset and approval can take up to 10 days.

9. Phase Status

Phase	Status	Notes
Phase 1 — Backbone Fine-Tuning	**✓ COMPLETE**	Checkpoint committed to Kaggle Output
Phase 2 — GAT Training	**IN PROGRESS**	Checkpoint available (gat_phase2_best.pt)
Phase 3 — Three-Head Training	*PENDING*	Emotion + MH heads trained; severity head pending DAIC-WOZ
Phase 4 — LLM Integration	*PENDING*	*Groq API integration — awaiting Phase 3 completion*